

Resource Allocation for Mini-batch SGD via Adaptive Batch Sizes

Scott Sievert



Agenda

- Motivation
- Results
- Practical implementation, with experiments

Joint work with
Zachary Charles



Motivation



Motivation: SGD

$$\min_{\boldsymbol{w} \in \mathbb{R}^d} F(\boldsymbol{w}) := \frac{1}{n} \sum_{i=1}^n f(\boldsymbol{w}; \boldsymbol{z}_i)$$

In deep learning, typically solved via SGD (or a variant):

$$\boldsymbol{w}_{k+1} = \boldsymbol{w}_k - \frac{\gamma}{B} \sum_{i=1}^B \nabla f(\boldsymbol{w}_k; \boldsymbol{z}_{i_j})$$

...with a static batch size B (another hyperparameter)

Extremely Large Minibatch SGD: Training ResNet-50 on ImageNet in 15 Minutes

[Takuya Akiba](#), [Shuji Suzuki](#), [Keisuke Fukuda](#)

(Submitted on 12 Nov 2017)

We demonstrate that training ResNet-50 on ImageNet for 90 epochs can be achieved **in 15 minutes with 1024 Tesla P100 GPUs**. This was made possible by using **a large minibatch size of 32k**. ...

Accurate, Large Minibatch SGD: Training ImageNet in 1 Hour

[Priya Goyal](#), [Piotr Dollár](#), [Ross Girshick](#), [Pieter Noordhuis](#), [Lukasz Wesolowski](#), [Aapo Kyrola](#), [Andrew Tulloch](#), [Yangqing Jia](#), [Kaiming He](#)

(Submitted on 8 Jun 2017 ([v1](#)), last revised 30 Apr 2018 (this version, v2))

... trains ResNet- 50 with **a minibatch size of 8192 on 256 GPUs in one hour**, ...

It'd be nice to train quickly
without having to buy 100's of GPUs

To do that, let's *grow* the batch size:

$$\boldsymbol{w}_{k+1} = \boldsymbol{w}_k - \frac{\gamma}{B_k} \sum_{i=1}^{B_k} \nabla f(\boldsymbol{w}_k; \boldsymbol{z}_{i_j})$$

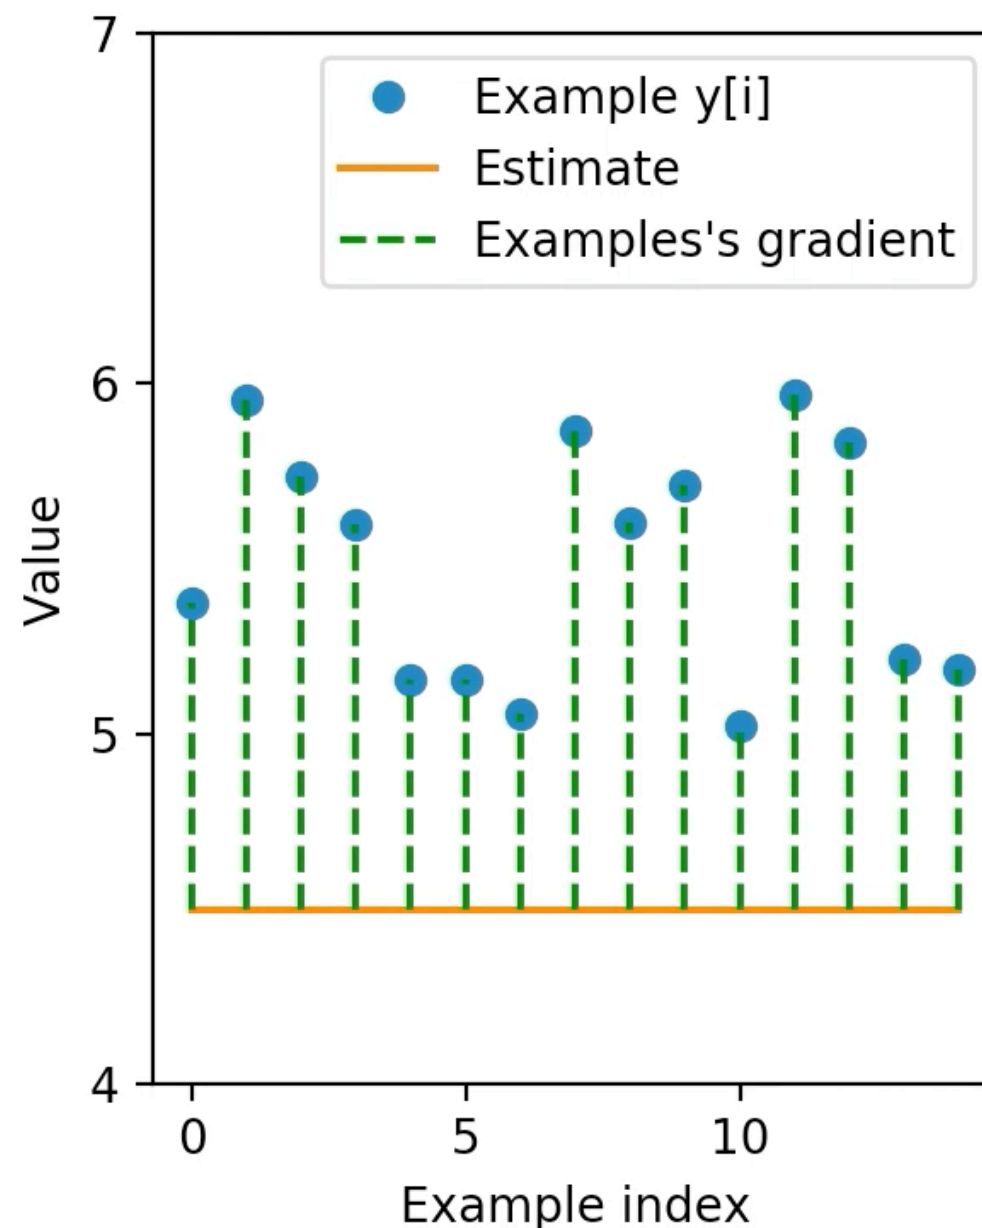
Why?

Intuition for batch size growth

Why have a large batch size with poor initialization?

Let's fit the simplest linear model:

$$\min_{x \in \mathbb{R}} \frac{1}{2} \sum_{i=1}^n (x - y_i)^2$$



Batch size growth

Let's grow the batch size via

$$B_k = \left\lceil \frac{c}{F(\mathbf{w}_k) - F^*} \right\rceil$$

a constant

best possible loss
(possibly unknown)

Poor initialization: $B_k = 1$

At optimum: $B_k = \infty$

Main result: summary

Feature of SGD: no dependence on
number of examples in training set

Feature of gradient descent (GD):
few model updates

When the batch size is $B_k = \left\lceil \frac{c}{F(\mathbf{w}_k) - F^*} \right\rceil$, then

the same number of

- model updates as GD is required
- examples as SGD is required

I have formal theorems backing this statement.

Main result

To achieve a model x with training loss $F(x_k) - F^* \in [\epsilon/2, \epsilon]$,

| Function class | SGD | Adaptive batch sizes | Gradient descent |
|-------------------|-----------------------------|---------------------------------|---------------------------------|
| Smooth and convex | $\mathcal{O}(1/\epsilon^2)$ | $\mathcal{O}(1/\epsilon)$ | $\mathcal{O}(1/\epsilon)$ |
| Strongly convex* | $\mathcal{O}(1/\epsilon)$ | $\mathcal{O}(\log(1/\epsilon))$ | $\mathcal{O}(\log(1/\epsilon))$ |

model updates need to be completed and

| Function class | SGD | Adaptive batch sizes | Gradient descent |
|-------------------|-----------------------------|--|-----------------------------------|
| Smooth and convex | $\mathcal{O}(1/\epsilon^2)$ | $\mathcal{O}(1/\epsilon^2)$ | $\mathcal{O}(n/\epsilon)$ |
| Strongly convex* | $\mathcal{O}(1/\epsilon)$ | $\mathcal{O}(\log(1/\epsilon)/\epsilon)$ | $\mathcal{O}(n \log(1/\epsilon))$ |

examples need to be processed

Similar results for reaching a saddle point with non-convex function

*actually a generalization of strongly convex [3]

1. Bubeck et al. *Convex optimization: Algorithms and complexity*. 2015.
2. Hamed Karimi et al. *Linear convergence of gradient and proximal-gradient methods under the polyak-łojasiewicz condition*. 2016.
3. Yurii Nesterov. *Introductory lectures on convex optimization: A basic course*. 2013.

These theorems are confirmed with simulations.

Limitations

- 1. Using the *entire* train dataset every model update takes a long time
- 2. How does the model perform on unseen data?
- 3. GPU memory is finite

Algorithm 1 Mini-batch SGD with dampening of noise in gradient approximation

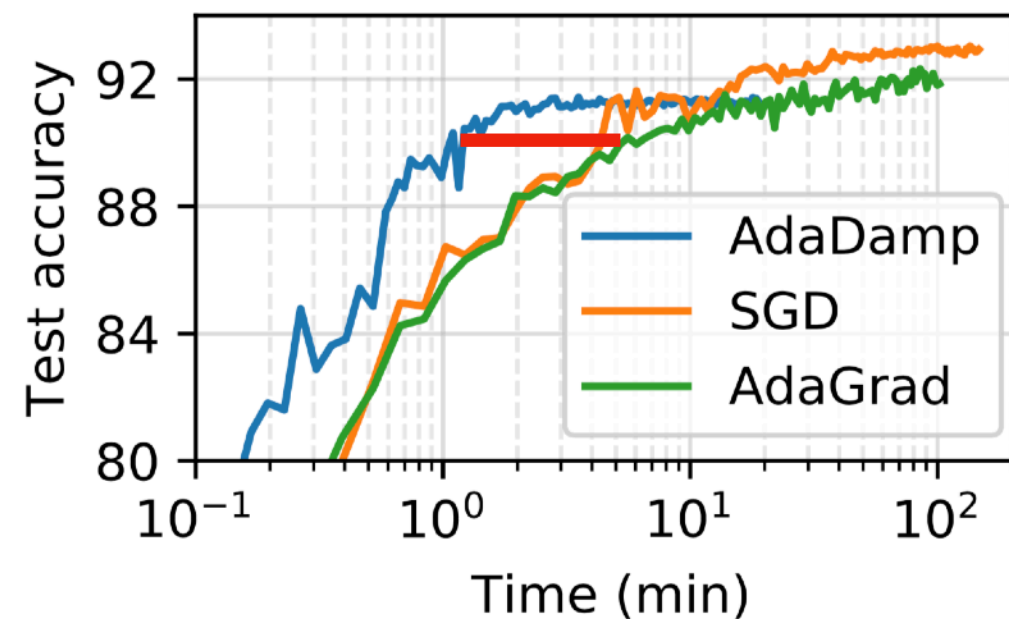
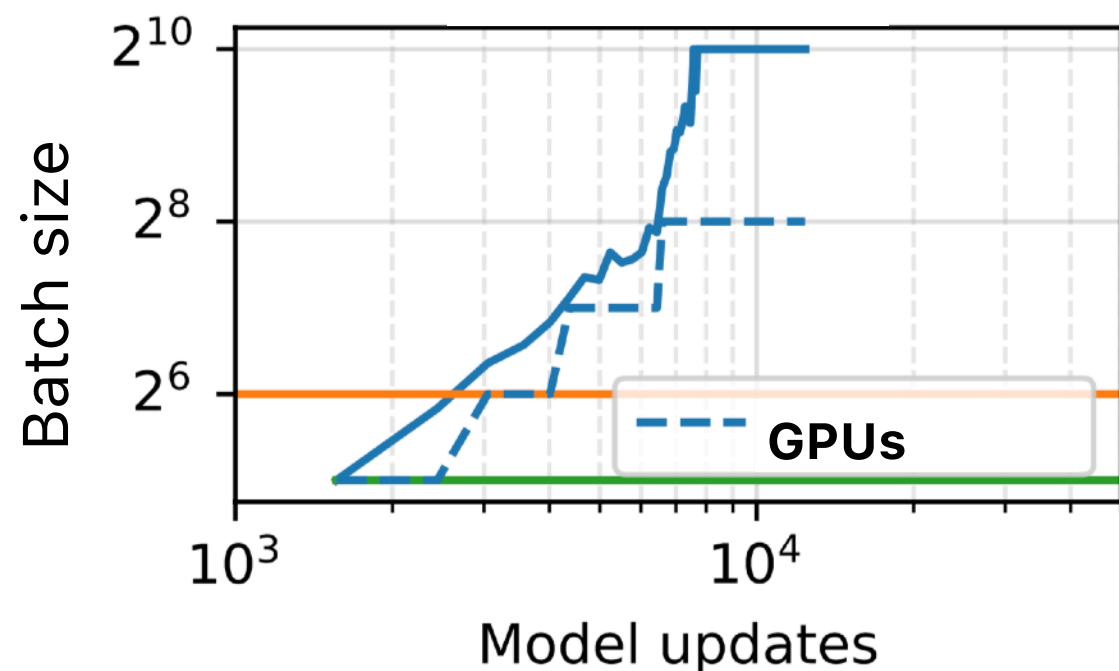
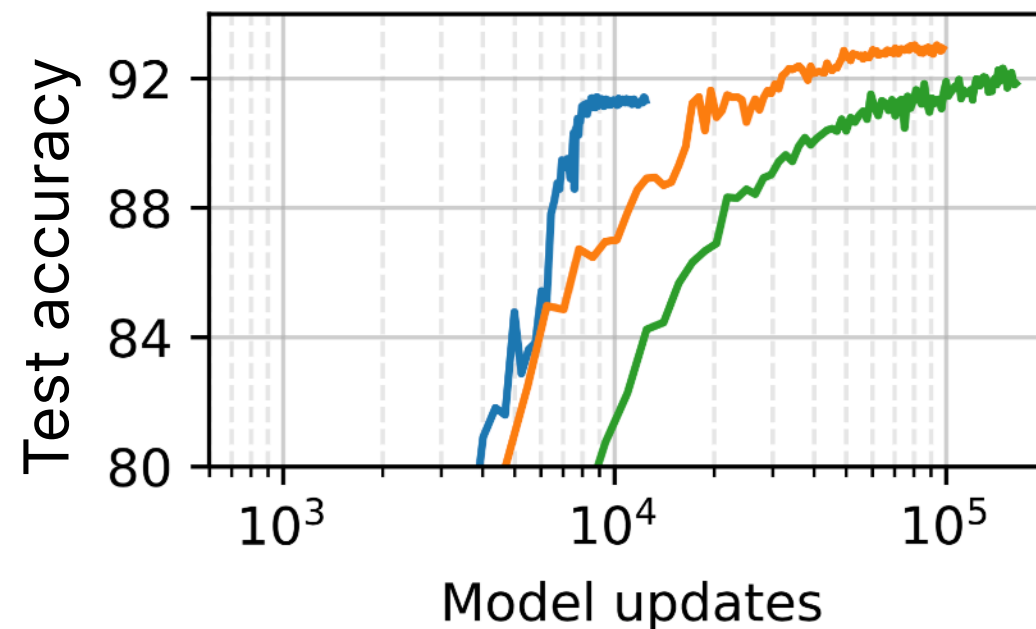
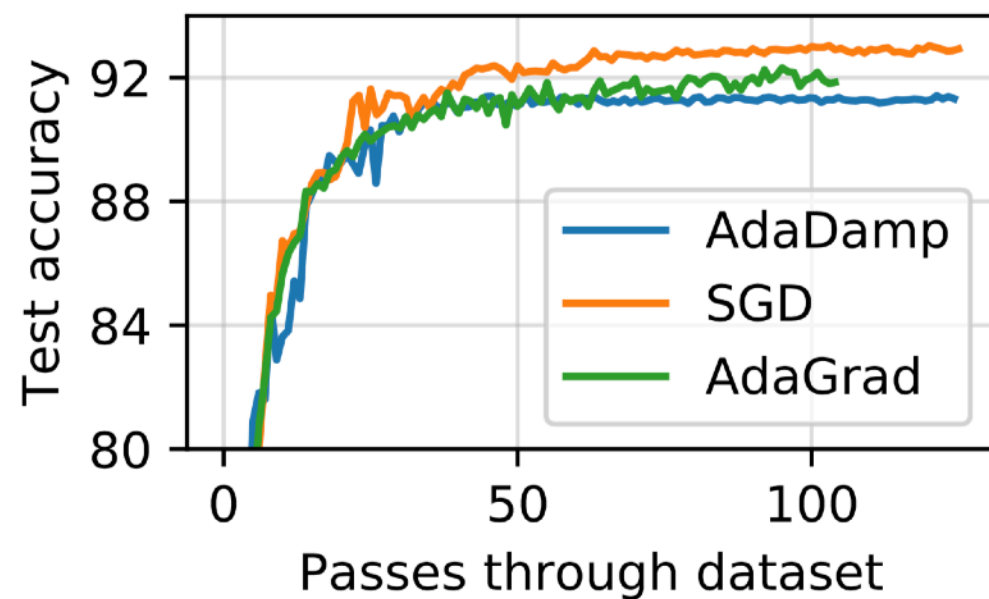
```
1: procedure ADADAMP(initial batch size  $B_0$ , initial model  $x_0$ , step size  $\gamma$ , maximum batch size  $B_{\max}$ )
2:   for  $k \in [0, 1, 2, \dots]$  do
3:     if  $k = 0$  then
4:        $c \leftarrow B_0(F(x_0) - F^*)$ 
5:        $B_k \leftarrow \lceil c / (F(x_k) - F^*) \rceil$ 
6:        $\gamma' \leftarrow \gamma$ 
7:       if  $B_k > B_{\max}$  then
8:          $\gamma' \leftarrow \gamma B_{\max} / B_k$ 
9:          $B_k = B_{\max}$ 
10:       $x_{k+1} \leftarrow \text{TRAIN}(x_k, \gamma', B_k)$ 
return  $x_k$ 
```

▷ TRAIN computes n gradients

Experiment

ResNet-34 + CIFAR-10

Brief hyperparameter tuning



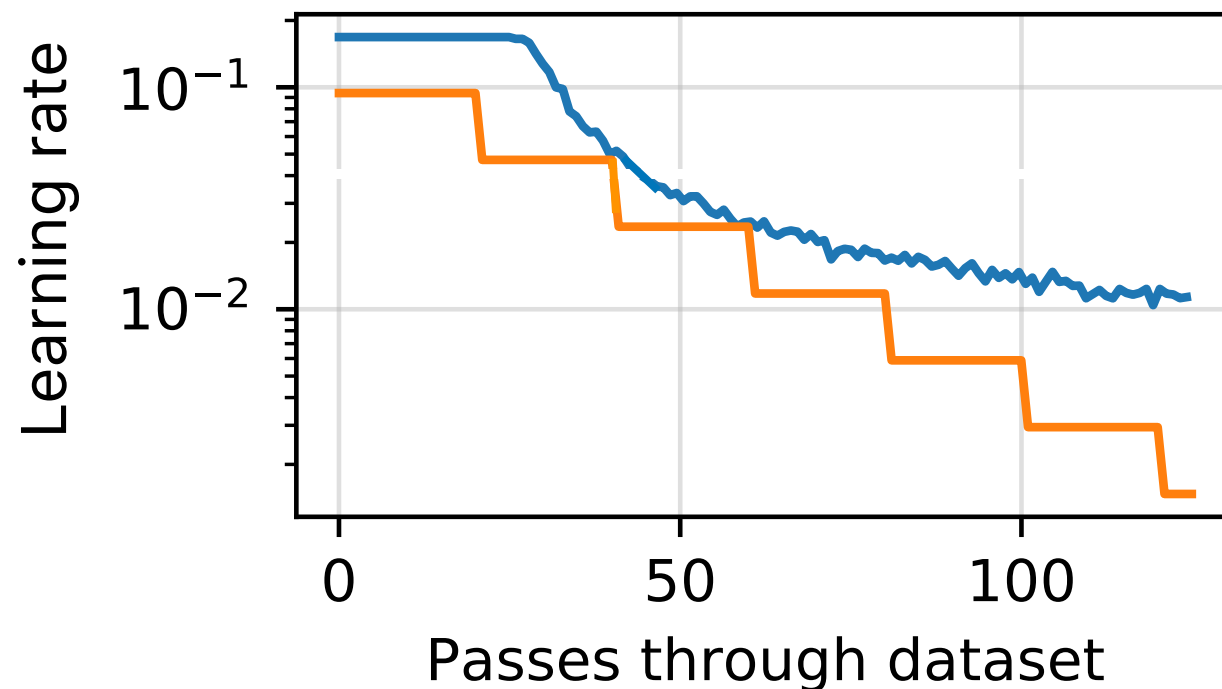
All plots assume oracle provides $F(x_k) - F^*$

All plots assume oracle provides $F(x_k) - F^*$

How can the batch size be estimated?

Our upper bounds suggest $B_{\text{epoch}} \propto r \cdot \text{epoch}$

Similar method uses $B_{\text{epoch}} \propto r^{\text{epoch}}$



This needs more investigation

Thanks!

Questions?

Future work

after this work is finished

How does this method generalize?

Bounds

Convex and smooth

$$F(\mathbf{x}_k) - F^* \leq \mathcal{O}(r^k)$$

α -PL
(generalization of
strongly convex)

$$F(\mathbf{x}_k) - F^* \leq \mathcal{O}\left(\frac{1}{k}\right)$$

Non-convex

$$\min_{k=0,\dots,T-1} \|\nabla F(\mathbf{x}_k)\|_2^2 \leq \mathcal{O}\left(\frac{1}{k}\right)$$

Max batch size?

Computer Science > Machine Learning


A Bayesian Perspective on Generalization and Stochastic Gradient Descent

Samuel L. Smith, Quoc V. Le

(Submitted on 17 Oct 2017 ([v1](#)), last revised 14 Feb 2018 (this version, v3))

... We also demonstrate that, when one holds the learning rate fixed, **there is an optimum batch size which maximizes the test set accuracy.** ...

Entire train dataset?



$$F(\mathbf{x}_k) - F^* \leq \mathcal{O}\left(\frac{1}{k}\right)$$

*For convex and non-convex functions. For convex, relies on

Rule of thumb can likely be developed